

6 High Availability

Sunday, 15 September 2019

9:39 PM

ELB / Route53

“Everything fails, all the time.”

High availability (HA) is about ensuring that your **downtime is minimized** as much as possible, without the need for human intervention.

Levels of Availability:

	Percent of Uptime	Max Downtime per Year	Equivalent Downtime per Year
1 Nine	90%	36.5 days	36.5 days
2 Nines	99%	3.65 days	3.65 days
3 Nines	99.9%	8.76 hrs	8.76 hrs
4 Nines	99.99%	52.6 min	52.6 min
5 Nines	99.999%	5.25 min	5.25 min

your application's
without the need for

**ivalent
ntime
Day**

2.4 hrs

14 min

86 sec

3.6 sec

86 sec

Avoid **SPOF**

Assume everything fails, and design

Recovery Time Objective (**RTO**)

How quickly must the system recover?

Recovery Point Objective (**RPO**)

How much data can you afford to lose?

How much money do you need to invest to meet these objectives?

Fault tolerance:

The **built-in redundancy** of an application's components.

Recoverability:

The process, policies, and procedures related to recovering from a catastrophic event.

Scalability:

The ability of an application to **accommodate growth**.

gn backwards.

vest to meet those

onents.

restoring service after a

with without changing

design.

HA on-premises **vs** HA on AWS

Ex: single RDS instance vs Multi-AZ

AWS services and high availability

- Inherently HA services
- HA with right architecture

ELB

3

lity

A managed load balancing service that **distributes traffic** across multiple Amazon EC2 instances.

Elastic Load Balancing:

- **Distributes load** between instances.
- Recognizes and responds to **unhealthy** instances.
- Can be **public** or **internal-facing**.
- Uses **HTTP**, **HTTPS**, and **TCP** protocols.
- Each load balancer is given a **public DNS** name.
 - **Internet-facing** load balancers have DNS names that publicly resolve to the **public** IP addresses of the load balancer's nodes.

Internal load balancers have DNS names that resolve to the **private** IP addresses of the load balancer's nodes.

DIY	Elastic Load Balancing
<ul style="list-style-type: none">• Un-managed (managed by you)• You create a load balancer on Amazon EC2 that fits your requirements• Scaling must be handled by you	<ul style="list-style-type: none">• Managed by AWS• Usually the most cost-effective solution• Automatically scales

s incoming application

instances.

S name.

S names which
s of the load

hich publicly resolve
cer's nodes.

Load Balancing

S

t cost-efficient

aled

Elastic Load Balancing provides the features:

- Health checks
- Cross-zone load balancing
- Proxy Protocol
- Sticky sessions
- Connection draining

Classic Load Balancer

Application Load Balancer

Network Load Balancer



ne following

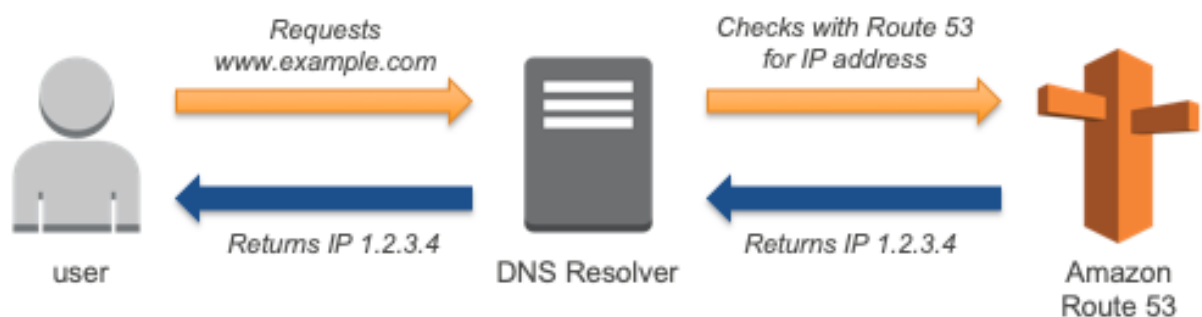
Improve the architecture exercise

se

EIP for HA

What about HA across regions?

Route 53 DNS Resolution



Reliable

- Redundant locations
- Backed with 100% Service Level Agreement (SLA)

Easy to Use

- Console
- Programmatic API
- Domain name management

Fast

- Worldwide anycast network
- Fast propagation of changes

Cost-Effective

- Inexpensive rates
- Pay-as-you-go model

Integrated with AWS

- ELB-Alias Queries
- Latency-based routing

Flexible

- Geolocation routing
- Weighted round robin
- Self-aliasing

Kind of Routing Route 53 supports

- **Simple routing:** Single server environment
- **Weighted round robin:** Assign weights to sets to specify the frequency.
- **Latency-based routing:** Helps to improve applications.
- **Health check and DNS failover:** Fail over if your primary site becomes unreachable.
- **Geolocation routing:** Specify **geographic** continent, by country, or by state in the U

ports?

nts.

o resource record

e your global

r to a backup site if

c locations by
United States.

Use Case: Multi-Region Deployment



Name	Type	Value
amgogreen.com	ALIAS	some-elb-name.us-west-2.elb.amazonaws.com
amgogreen.com	ALIAS	some-elb-name.ap-southeast-2.elb.amazonaws.com

Typical Architecture

Record Sets
CNAME www

elastic_load_balancer
Routing Policy = Failover
Record Type = Primary

Amazon S3 website
Routing Policy = Failover
Record Type = Secondary





east-2.elb.amazonaws.com



naws.com

mazonaws.com

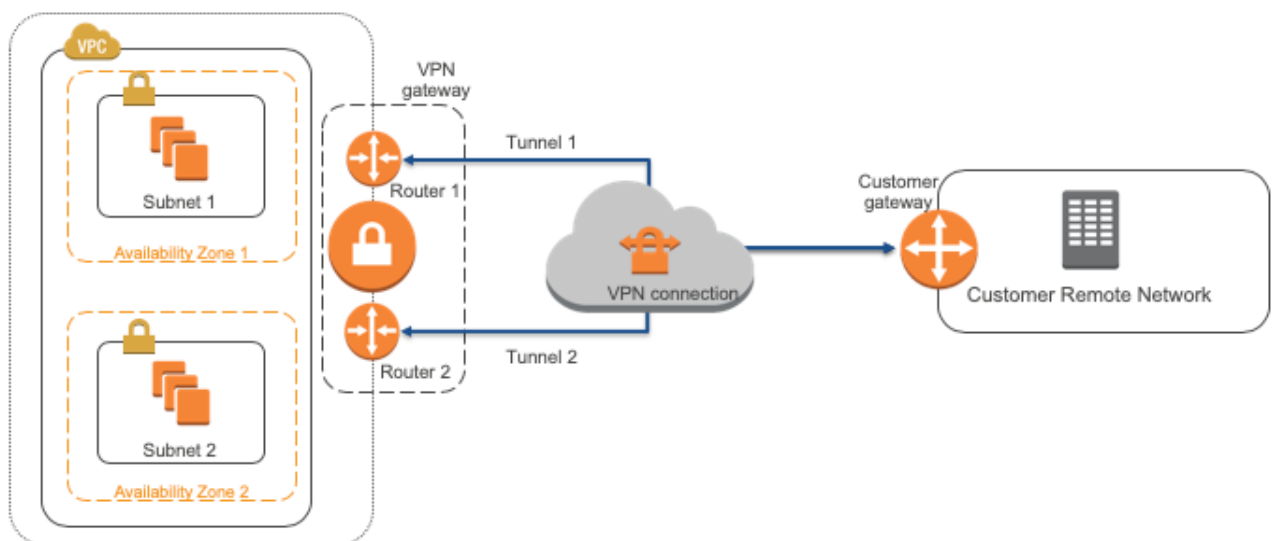




- Create a domain that uses Amazon Route 53 as the DNS service: see <http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/>
- Migrate an existing domain to Amazon Route 53: see <http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/MigratingDNS.html>
- Create a subdomain that uses Amazon Route 53 without migrating the parent domain: see <http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/CreatingNewSubdomain.html>
- Migrate a subdomain to Amazon Route 53 without migrating the parent domain: see <http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/MigratingSubdomain.html>

Create alias record sets for Elastic Load Balancing:

<http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/resource-record-sets.html#resource-record-sets-alias>

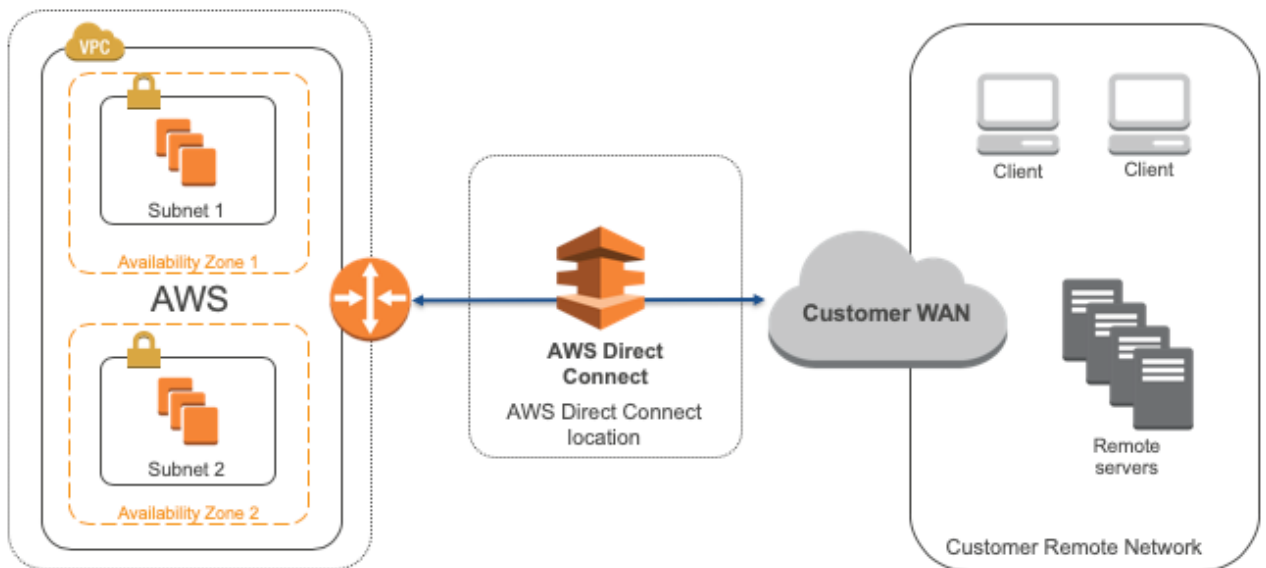



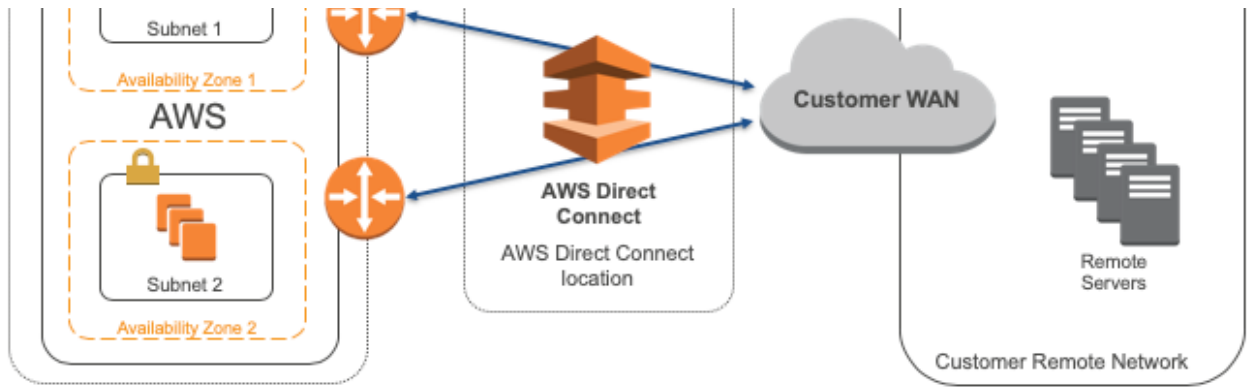


main: see
[main.html](#)

see
[in.html](#)

[ts-choosing-alias-non-alias.html](#)





Lab :

1. *Inspect the resources created automatically for you*
2. *Start your web server's PHP application*
3. *Create an AMI of your web server.*
4. *Create a second public subnet in a second AZ and a server in that subnet.*
5. *Create a load balancer and attach it to both subnets.*
6. *Create two private subnets with two instances in each.*
7. *Create a NAT Gateway and route traffic from private subnets to the Internet.*
8. *Test the private instances' connections to the NA*



our lab.

and launch a new web

ets.

side of them.

ate instances to it.

T Gateway.

